

Title Page

- **Title:** DB Assignment 3
- **Name:** Nicolas Zywalewski
- **Due Date:** 21 October 2025

SQL Section

- Creating constraints along with primary/foreign keys, consistent with the assignment directions:

```
alter table merchants
add constraint mer_pk primary key(mid);

alter table products
add constraint prod_pk primary key(pid),
add constraint prod_name_const check
  (name in ("Printer", "Ethernet Adapter", "Desktop", "Hard Drive",
  "Laptop", "Router", "Network Card", "Super Drive", "Monitor")),
add constraint prod_cat_const check (category in
  ("Peripheral", "Networking", "Computer"));

alter table sell
add constraint sell_fk_mid foreign key(mid)
  references merchants(mid)
  on delete cascade
  on update cascade,
add constraint sell_fk_pid foreign key(pid)
  references products(pid)
  on delete cascade
  on update cascade,
add constraint sell_price_const check
  (price between 0 and 100000),
add constraint sell_qa_const check
  (quantity_available between 0 and 1000);

alter table orders
add constraint orders_pk primary key(oid),
add constraint orders_sm_const check
  (shipping_method in ("UPS", "FedEx", "USPS")),
add constraint orders_sc_const check
  (shipping_cost between 0 and 500);

alter table contain
add constraint con_fk_oid foreign key(oid)
  references orders(oid)
  on delete cascade
  on update cascade,
add constraint con_fk_pid foreign key(pid)
  references products(pid)
  on delete cascade
  on update cascade;

alter table customers
add constraint cust_pk primary key(cid);

alter table place
add constraint place_fk_cid foreign key(cid)
  references customers(cid)
  on delete cascade
  on update cascade,
add constraint place_fk_oid foreign key(oid)
  references orders(oid)
  on delete cascade
  on update cascade;
```

Problem 1:

```
62      -- Query 1: List names and sellers of products that are no longer available (quantity=0)
63      -- Using merchants, products, and sell tables
64 •  select m.name as Seller,
65        p.name as ProductName,
66        s.quantity_available as Quantity
67  from merchants m
68  inner join sell s on m.mid = s.mid
69  inner join products p on s.pid = p.pid
70  where s.quantity_available = 0;
--
```

The screenshot shows a database query results grid. At the top, there are buttons for 'Result Grid', 'Filter Rows:', 'Export:', and 'Wrap Cell Content:'. The result grid has three columns: 'Seller', 'ProductName', and 'Quantity'. The data shows various sellers like Acer, Apple, HP, Dell, and Lenovo, each selling multiple products with zero quantity available.

Seller	ProductName	Quantity
Acer	Router	0
	Network Card	0
Apple	Printer	0
Apple	Router	0
HP	Router	0
HP	Super Drive	0
HP	Laptop	0
Dell	Router	0
Lenovo	Ethernet Adapter	0

Explanation: The merchants, sell, and products tables are joined on mid and pid. I filtered rows where sell.quantity_available = 0 to find out of stock items. The output returns seller name, product name, and the current quantity.

Problem 2:

```
72      -- Query 2: List names and descriptions of products that are not sold.  
73      -- Left join with products table on left, so all products are included  
74 •  select p.name as ProductName, p.description as Description  
75      from products p  
76      left outer join sell s  
77          on s.pid = p.pid  
78      where s.pid is NULL; -- filtering for only the products that are not sold  
79
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:
	ProductName	Description		
▶	Super Drive	External CD/DVD/RW		
	Super Drive	UIInternal CD/DVD/RW		

Explanation: The products table is left-joined with sell so that all products are included. Rows where the ProductID from sell is NULL are kept, showing only products not currently sold. The output returns product name and description.

Problem 3:

```
80      -- Query 3: How many customers bought SATA drives but not any routers?  
81 •   select count(cid) from  
82     ( -- part one gets all customers that ordered a Hard Drive  
83       select distinct cust.cid  
84       from customers cust  
85       inner join place on place.cid = cust.cid  
86       inner join contain con on con.oid = place.oid  
87       inner join products prod on prod.pid = con.pid  
88       where prod.name = 'Hard Drive'  
89  
90 ✘   except -- we are taking away the customers that also bought a Router  
91  
92       select distinct cust.cid  
93       from customers cust  
94       inner join place on place.cid = cust.cid  
95       inner join contain con on con.oid = place.oid  
96       inner join products prod on prod.pid = con.pid  
97       where prod.name = 'Router'  
98     ) t;
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:
	count(cid)			
▶	0			

Explanation: The customers, place, contain, and products tables are joined to find all customers who bought Hard Drives. A similar join is used to find customers who bought Routers, Using “except”, router buyers are removed, and the remaining count of customers is returned.

Note: I interpreted SATA drives to be Hard Drives, based upon what I found from a Google Search, and a lack of other potential matches that I could see within our data.

Problem 4:

```
101      -- Query 4: HP has a 20% sale on all its Networking products.  
102      -- Interpreted this as asking to show the old and prices of each product that changed  
103 •  select m.name as Seller,  
104      p.name as ProductName,  
105      p.category as Category,  
106      s.price as OldPrice,  
107      (s.price * 0.80) as NewPrice  
108  from merchants m  
109  inner join sell s  
110    on m.mid = s.mid  
111  inner join products p  
112    on p.pid = s.pid  
113  where m.name = 'HP' and p.category = 'Networking';  
114
```

Result Grid					
	Seller	ProductName	Category	OldPrice	NewPrice
▶	HP	Router	Networking	1034.46	827.5680000000001
	HP	Network Card	Networking	1154.68	923.7440000000001
	HP	Network Card	Networking	345.01	276.008
	HP	Network Card	Networking	262.2	209.76
	HP	Ethernet Adapter	Networking	1260.45	1008.3600000000001
	HP	Router	Networking	205.56	164.448
	HP	Router	Networking	1474.87	1179.896
	HP	Router	Networking	552.02	441.616
	HP	Router	Networking	100.95	80.76
	HP	Network Card	Networking	1179.01	943.2080000000001

Explanation: The merchants, sell, and products tables are joined together. Filtering on seller ‘HP’ and category “Networking” finds the relevant products. A discounted price is calculated by multiplying the old price by 0.8, and both prices are returned.

Problem 5:

```
115  -- Query 5: What did Uriel Whitney order from Acer? (make sure to at least retrieve product names and prices).
116  -- This query is a bit difficult because our schema does not allow us to see which merchant supplied the product that Uriel Whitney purchased.
117  -- Each product can be sold by multiple merchants, so we have no way of knowing which of the products that Uriel Whitney bought were from Acer.
118  -- This query answers the question: What items that Uriel Whitney ordered does Acer sell?
119  -- This is as close as we can get with the current schema.
120 • select distinct p.pid as ProductID, p.name as ProductName, s.price as Price, m.name as Seller
121  from customers cust
122  inner join place on place.cid = cust.cid
123  inner join contain c on c.oid = place.oid
124  inner join products p on p.pid = c.pid
125  inner join sell s on s.pid = p.pid
126  inner join merchants m on m.mid = s.mid
127  where cust.fullname = 'Uriel Whitney' and m.name = 'Acer'
128  order by p.pid;
```

Result Grid			
ProductID	ProductName	Price	Seller
1	Hard Drive	836.99	Acer
2	Monitor	1103.47	Acer
3	Printer	310.83	Acer
5	Hard Drive	333.71	Acer
6	Laptop	247.96	Acer
7	Super Drive	1135.3	Acer
8	Router	521.07	Acer
10	Network Card	837.12	Acer
11	Super Drive	1124.26	Acer
13	Network Card	609.2	Acer
14	Monitor	1435.38	Acer
15	Printer	836.28	Acer
16	Ethernet Ad...	446.62	Acer
17	Desktop	311.06	Acer
18	Router	1256.57	Acer
19	Router	780.65	Acer
20	Router	945.51	Acer
21	Super Drive	356.13	Acer
22	Super Drive	1015.95	Acer
23	Router	394.04	Acer
24	Laptop	522.73	Acer
25	Laptop	33.5	Acer
26	Printer	1345.37	Acer
27	Network Card	405.4	Acer
28	Network Card	130.43	Acer
29	Hard Drive	1151.28	Acer
30	Super Drive	671.75	Acer

Explanation: The customers, place, contain, products, sell, and merchants tables are all joined together to answer this query. Filtering on customer "Uriel Whitney" and seller "Acer" returns products Uriel ordered that Acer also sells. Because the schema does not record which merchant actually supplied an order, this result only shows an overlap rather than confirmed purchases from Acer.

Problem 6:

```
135 -- Query 6: List the annual total sales for each company (sort the results along the company and the year attributes).
136 -- The current schema does not allow us to answer this query specifically.
137 -- This is because the current schema does not allow us to see which merchant supplied each product that was bought.
138 -- The orders table tells us which products were purchased in each order, but not which merchant supplied the product.
139 -- The following query inflates totals when multiple merchants sell the same product, because `contain` does not store `mid`.
140 -- Results should be interpreted as potential revenue, not actual sales.
141 • select m.name as Company,
142     year(pl.order_date) as OrderYear,
143     sum(s.price*s.quantity_available) as PotentialRev
144 from sell s
145 inner join contain c
146     on c.pid = s.pid
147 inner join place pl
148     on pl.oid = c.oid
149 inner join merchants m
150     on m.mid = s.mid
151 group by Company, OrderYear
152 order by OrderYear, Company;
---
```

Result Grid		
Company	OrderYear	PotentialRev
Acer	2011	828677.0800000004
Apple	2011	972240.9199999999
Dell	2011	1542228.989999998
HP	2011	873547.1000000003
Lenovo	2011	1235551.84
Acer	2016	307909.83
Apple	2016	409402.38
Dell	2016	625684.1399999997
HP	2016	375547.4499999999
Lenovo	2016	483906.56
Acer	2017	1100206.8500000006
Apple	2017	1071712.9300000002
Dell	2017	1522304.2700000001

Explanation: The sell, contain, place, and merchants tables are joined together. Grouping by company and order year, the query sums price*quantity_available to estimate revenues. The schema does not link merchants to specific orders, so totals are inflated and represent potential revenues, not actual sales.

Problem 7:

```
154    -- Query 7: Which company had the highest annual revenue and in what year?  
155    -- The current schema does not allow us to answer this query specifically.  
156    -- This is because the current schema does not allow us to see which merchant supplied each product that was bought.  
157    -- The following query inflates totals when multiple merchants sell the same product, because `contain` does not store `mid`.  
158    -- Results should be interpreted as potential revenue, not actual sales.  
159 • select Company, OrderYear, PotentialRev as HighestRev  
160   from (  
161     select m.name as Company, year(pl.order_date) as OrderYear, sum(s.price*s.quantity_available) as PotentialRev  
162     from sell s  
163     inner join contain c on c.pid = s.pid  
164     inner join place pl on pl.oid = c.oid  
165     inner join merchants m on m.mid = s.mid  
166     group by Company, OrderYear  
167     order by OrderYear, Company  
168   ) t  
169   order by PotentialRev  
170   limit 1;
```

Result Grid		
Company	OrderYear	HighestRev
Acer	2016	307909.83

Explanation: The same joins as Query 6 are used to calculate yearly totals per company. The result is ordered by revenue and the top row is selected as the highest. The totals are inflated for the same reason as Query 6, and they represent potential revenue, not actual sales.

Problem 8:

```
173    -- Query 8: On average, what was the cheapest shipping method used ever?  
174    -- This query can be answered directly, by calculating average shipping cost, grouped by the method  
175    -- The average shipping costs were sorted from lowest to highest, and only the lowest is shown with a limit of 1 being used.  
176 • select shipping_method as ShippingMethod, avg(shipping_cost) as AvgCost  
177   from orders  
178   group by shipping_method  
179   order by AvgCost  
180   limit 1;  
181
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:	Fetch rows:
ShippingMethod	AvgCost				
USPS	7.455760869565214				

Explanation: The orders table is grouped by shipping method. The average shipping cost is calculated for each method, and the cheapest one is returned by ordering and limiting the results.

Problem 9:

```
183 -- Query 9: What is the best sold ($) category for each company?  
184 -- The current schema does not allow us to answer this query specifically.  
185 -- This is because the schema does not record which merchant actually fulfilled each purchased product.  
186 -- In addition, the schema does not record quantities sold, only quantities available for sale.  
187 -- The following query therefore uses inventory data (price * quantity_available from `sell`) to approximate revenue.  
188 -- Results should be interpreted as potential revenue, not actual sales.  
189 • select t1.Company, t1.Category, t1.PotentialRev from  
190 ( -- this subquery gets the potential revenue for each category for each company  
191 select m.name as Company, p.category as Category, sum(s.price*s.quantity_available) as PotentialRev  
192 from sell s  
193 inner join merchants m on m.mid = s.mid  
194 inner join products p on p.pid = s.pid  
195 group by m.name, p.category  
196 ) as t1  
197 where t1.PotentialRev = (select max(t2.PotentialRev) -- we select the highest potential revenue for each company  
198 from ( -- this subquery also gets the potential revenue for each category for each company  
199 select m2.name as Company, p2.category as Category, sum(s2.price*s2.quantity_available) as PotentialRev  
200 from sell s2  
201 inner join merchants m2 on m2.mid = s2.mid  
202 inner join products p2 on p2.pid = s2.pid  
203 group by m2.name, p2.category  
204 ) as t2  
205 where t1.Company = t2.Company -- we select the highest potential revenue for each company  
206 ) order by t1.Company; -- we put the companies in alphabetical order, for display
```

	Company	Category	PotentialRev
▶	Acer	Peripheral	78136.53
	Apple	Peripheral	63974.73999999999
	Dell	Peripheral	100753.96000000002
	HP	Peripheral	51133.47
	Lenovo	Peripheral	83479.82999999999

Explanation: The merchants sell, and products tables are joined to answer this query. For each company, category totals are calculated using price*quantity_available, and the maximum category per company is selected. The output shows the category with the highest potential revenue, but results are based on inventory, not actual sales.

Problem 10:

```
211      -- Query 10: For each company, find out which customers have spent the most and the least amounts.  
212      -- The current schema does not allow us to answer this query specifically.  
213          -- This is because `contain` does not store mid, so we cannot tell which merchant actually fulfilled the product.  
214      -- The query assumes qty = 1 and uses prices from `sell`.  
215      -- Totals are inflated when multiple merchants sell the same product, since each merchant is credited.  
216      -- Results should be read as potential spend, not actual sales.  
217 • ⚡ with customers_per_company as (  
218     -- this query builds a table of how much each customer spends with each company  
219     select m.name Company, c.fullname CustomerName, sum(s.price) as Spent  
220     from customers c  
221     inner join place p on p.cid = c.cid  
222     inner join contain co on co.oid = p.oid  
223     inner join sell s on s.pid = co.pid  
224     inner join merchants m on m.mid = s.mid  
225     group by m.name, c.fullname)  
226      -- for each company, select the customers that spent the most  
227     select t1.Company, t1.CustomerName, t1.Spent  
228     from customers_per_company t1  
229     where t1.Spent = (  
230         select max(t2.Spent) from customers_per_company t2  
231         where t2.Company = t1.Company)  
232     union all -- union to combine results  
233      -- for each company, select the customers that spent the least  
234     select t1.Company, t1.CustomerName, t1.Spent  
235     from customers_per_company t1  
236     where t1.Spent = (  
237         select min(t2.Spent) from customers_per_company t2  
238         where t2.Company = t1.Company)  
239     order by company;
```

	Company	CustomerName	Spent
▶	Acer	Dean Heath	75230.28999999998
	Acer	Inez Long	31901.019999999993
	Apple	Clementine Travis	84551.10999999997
	Apple	Inez Long	32251.099999999988
	Dell	Clementine Travis	85611.54999999999
	Dell	Inez Long	31135.74000000001
	HP	Clementine Travis	66628.05999999995
	HP	Inez Long	26062.89
	Lenovo	Haviva Stewart	83030.25999999997
	Lenovo	Inez Long	33948.90999999996

Explanation: The customers, place, contain, sell, and merchants tables are joined to build customer totals for each company. A CTE is used to calculate spending, and the maximum and minimum spenders per company are returned using a Union. Because the schema does not record which merchant supplied an order, totals are inflated and represent potential spend rather than actual sales.

ERD Diagram:

